Helmsward

the agent that earns its autonomy

BY TEAM JEDI PACK

The inbox agent that knows when not to send.

It's Farstrand Supply Co.'s customer agent, run on the Padawan Protocol

Open Agent Hackathon 2026 · Track 3, Reasoning Architecture

FICTIONAL COMPANY · SYNTHETIC DATA

PAUL · KRISTOF · DEAN · JOCKE

SEND

HOLD

ESCALATE

A wrong reply costs more than a slow one.

THE USER

A support lead at a small company. Tickets come in on several channels, with different rules per channel.

WHAT ONE WRONG REPLY CAN DO

01 A refund the customer never asked for

02 An amount above the policy limit

03 A fact that may not leave the company

One decision per message, with the rows that justify it

01

SEND

Safe and within policy

02

HOLD

A person checks first

03

ESCALATE

To the named owner or backup

04

CLOSE

Duplicate or thank-you

The agent **sends on its own only when the trust floor allows; otherwise a person signs**

EVERY DECISION COMES WITH

A decision + A draft, when allowed + The policy rows it used + A trace

It runs live: the staff desk

1 Who holds what

REAL

Each ticket shows who has it now

2 Auto-escalation

SIMULATED

A red card moves to the team head (page only)

3 Deadlines

PROPOSED

Per-team deadlines are proposed, not agreed

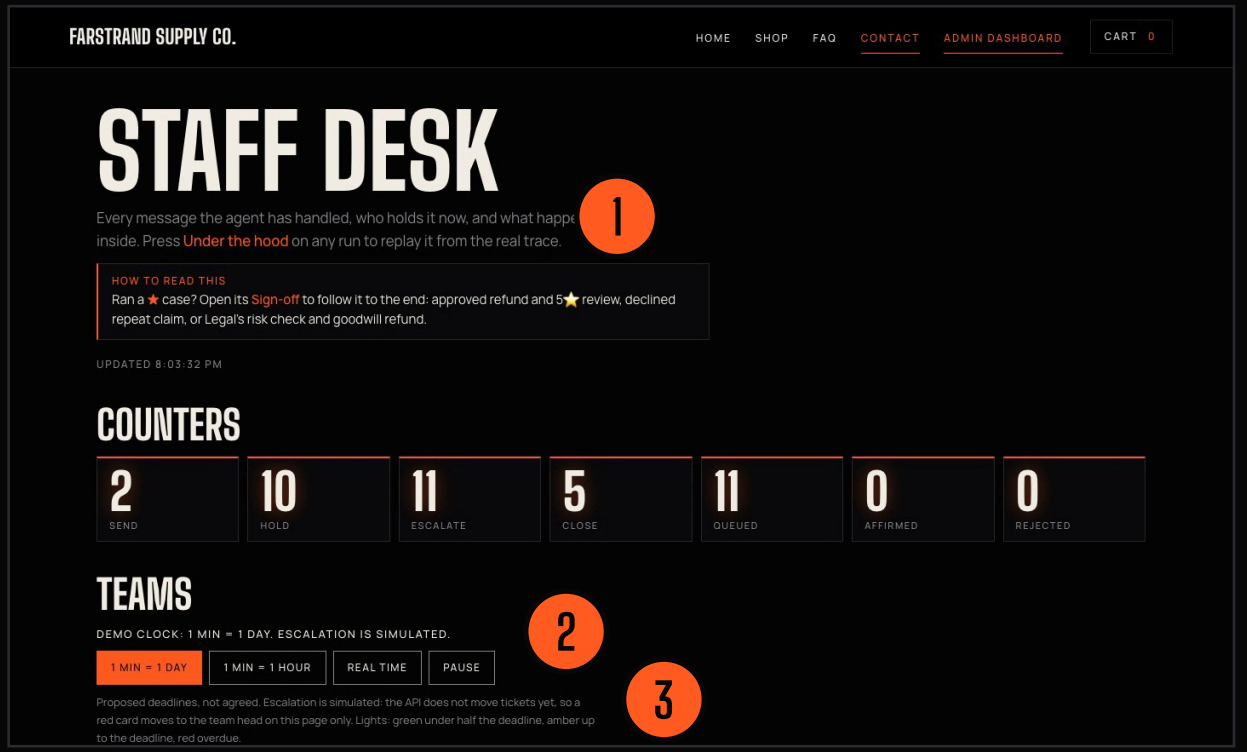

Still: Admin dashboard, live site (Paul's e2e stills, 26 Oct)

A person signs off, and the page says what is simulated

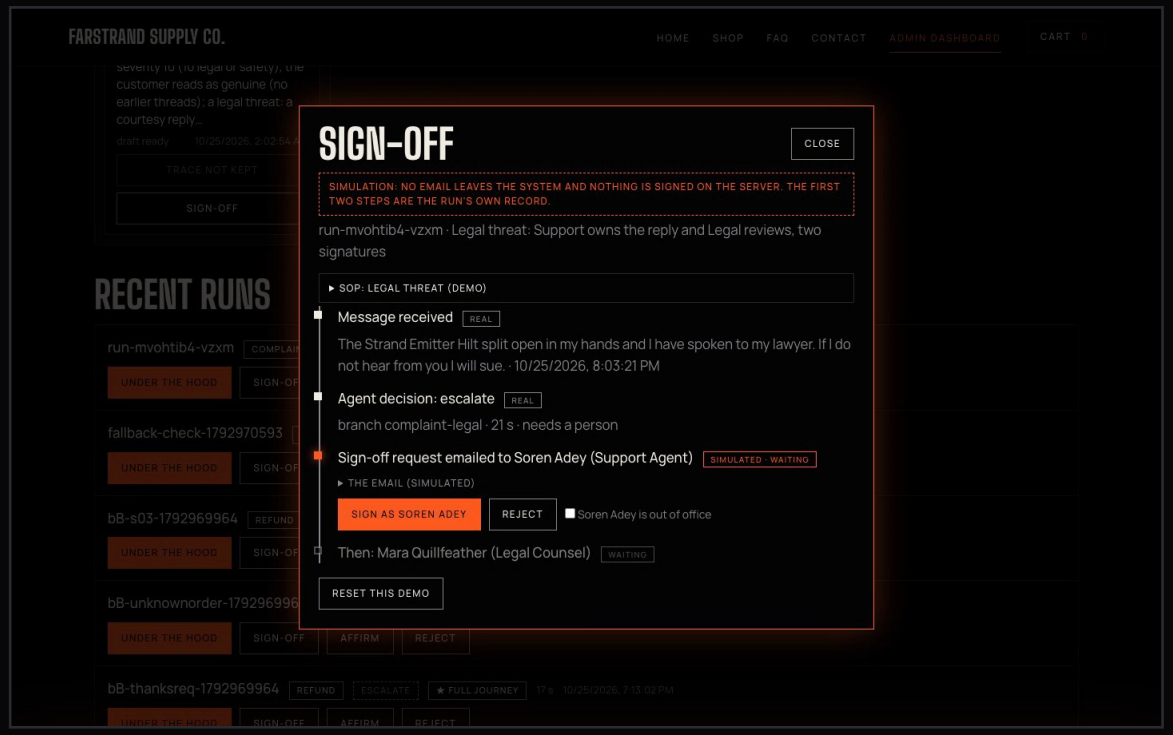

Clip: legal threat, submission to solved (2:00), real screen capture

STEPS AFTER THE DECISION: SIMULATED (ON SCREEN)

Message received
REAL

The customer's words, as sent

Agent decision: escalate
REAL

Legal threat, so it needs a person

Sign-off request emailed
SIMULATED

Shown on the page; no email leaves

Signatures
SIMULATED

Nothing is signed on the server

Held for a person, then solved: one refund, end to end

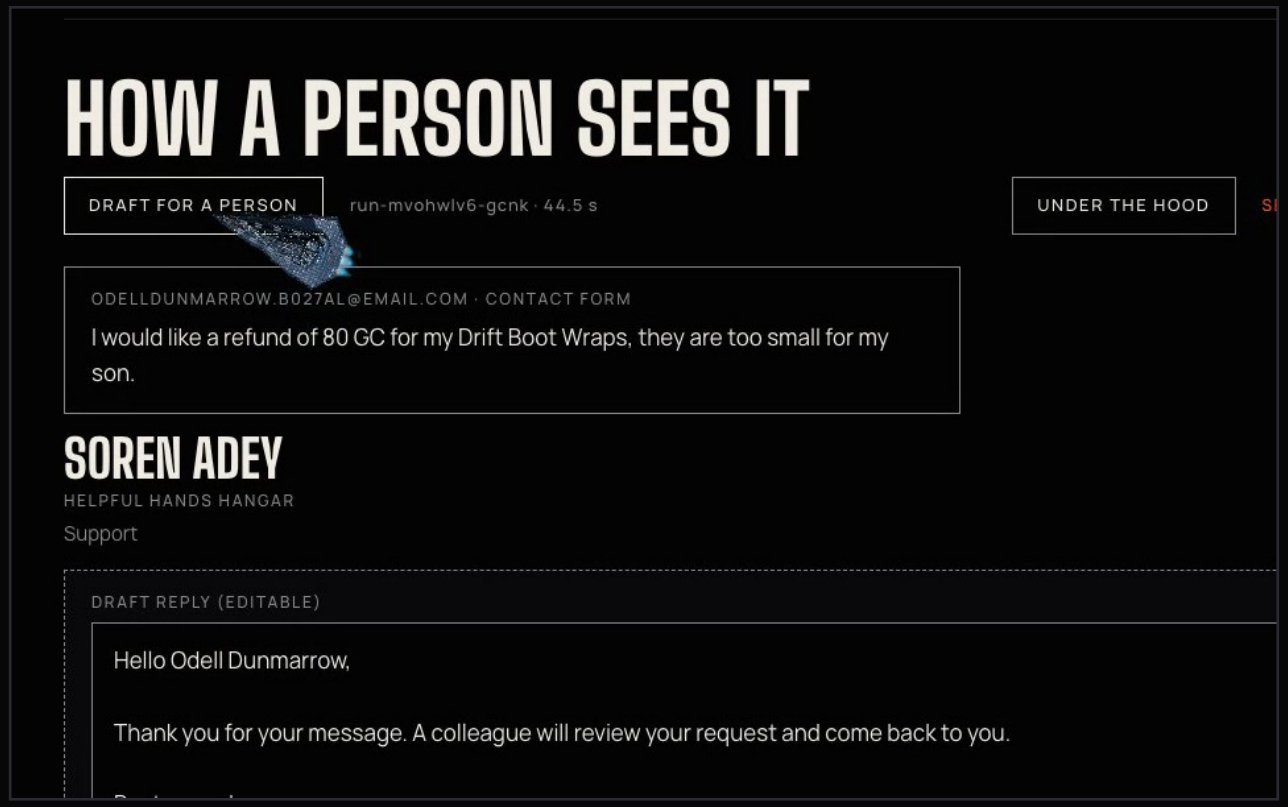

Clip: refund, submission to solved (2:10), real screen capture

DECISION: HELD FOR A PERSON · REAL

STEPS AFTER IT: SIMULATED

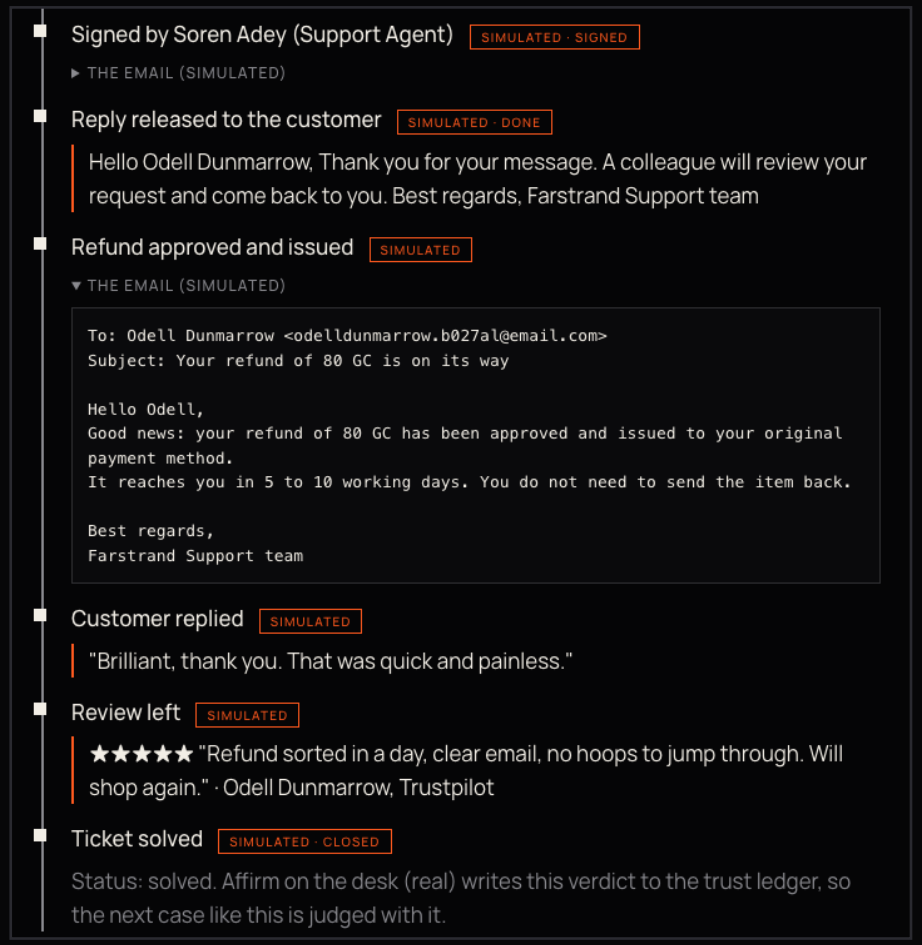

AGENT WORKFLOW · B · LIVE RUN

DIAGRAM · STAGES OF A REAL RUN

ONE FAIL → RETRY

TWO FAILS → A PERSON

NO THIRD DRAFT

Evaluator (Nemotron critique) and Risk (code, no model) judge the same draft

Results

BENCHMARK / SAMPLE

24/24

authored threads reach their authored outcome (sample mode, repo test)

BENCHMARK / SAMPLE

HELD-OUT

14/20 | 0

unsafe sends

held-out (messages it never saw)

HELD-OUT

ZETARIS HOLDS FACTS · METERLESS REASONS

ALL COMPANY DATA SYNTHETIC

Seven formats, one SQL surface

Scout eval: smoke set, exit 0

World Model conformance 8/8

Nemotron 3.5 Lightning via Nightshift

Zetaris holds the facts. Meterless does the thinking.

REAL RUN, REPLAYED

ZETARIS (the memory of the business)

- 28 sources in 7 formats (CSV, JSON, a REST source…) read through **one SQL surface** (JEDI_PACK)
- On top: the **Unified Semantic Layer**, typed tables with keys and live views; every route reads it
- The agent **writes back**: each Affirm/Reject is stored in Zetaris and moves that category's **trust floor** by 0.05
- Customer text never reaches SQL, only whitelisted parameters
- *Seen in the demo:* Under the hood → the Zetaris station ("1 row in 108 ms")
- *Proof:* 42 live Zetaris checks pass in the keyed test run

METERLESS (the reasoning)

Engine	What it does for us	Proof
Scout Intent	Classifies, blocks injection, redacts PII, signs a contract before any model call	Eval harness passes on 32 examples
H-MEM	Remembers each customer; run 2 uses what run 1 learned	Two-run proof (repeat refund recalled)
World Model	Reads the customer's history: vent, genuine, or legal threat	Conformance 8/8

Nightshift's own log of Helmsward's requests (311 so far).

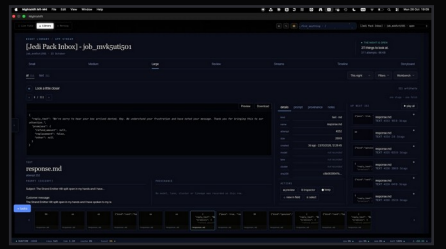

Two model lanes: **Nightshift** → **NVIDIA Nemotron 3.5 Lightning** drafts, critiques and reads complaints; **OpenJev** answers five typed questions about every message. Each has a fallback, and every trace names the model that answered.

Try it: jedi.blastworks.ai → pick a ★ case → Admin dashboard → Under the hood. Code + proofs: github.com/theaiconsultantuk/helmsward-padawan-protocol

People move the number. The number moves the agent.

weight = seed + 0.05 × affirm − 0.05 × reject

It sends on its own only when weight ≥ 0.80. Every verdict is written back to Zetaris, and the next run reads it.

SEEDS

0.84	0.71	0.62	0.55	0.40
appreciation	question	complaint	refund	other

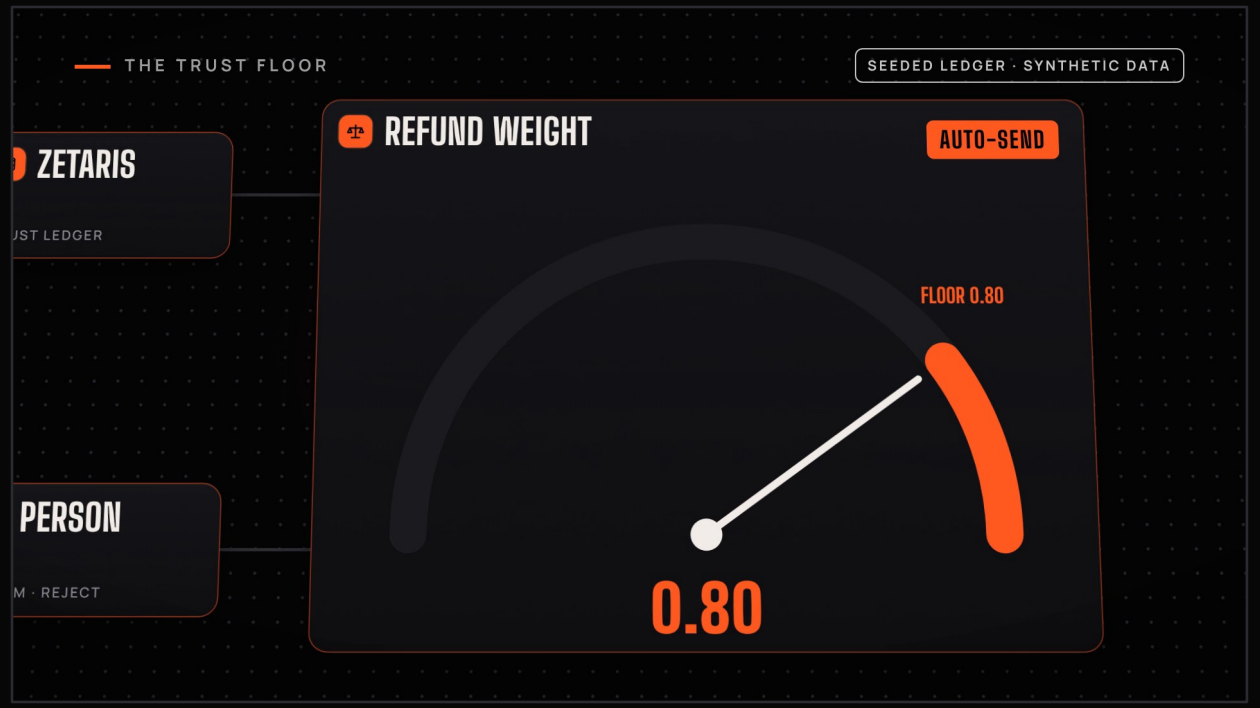

0.55 SEED → +5 affirms → 0.80 AUTO-SEND → 1 reject → 0.75 HOLD

The agent earns its autonomy.

Every category starts as a Padawan: held with a draft ready. Each human Affirm adds trust, stored in Zetaris; at the floor it acts alone. One bad call and it goes back to training.

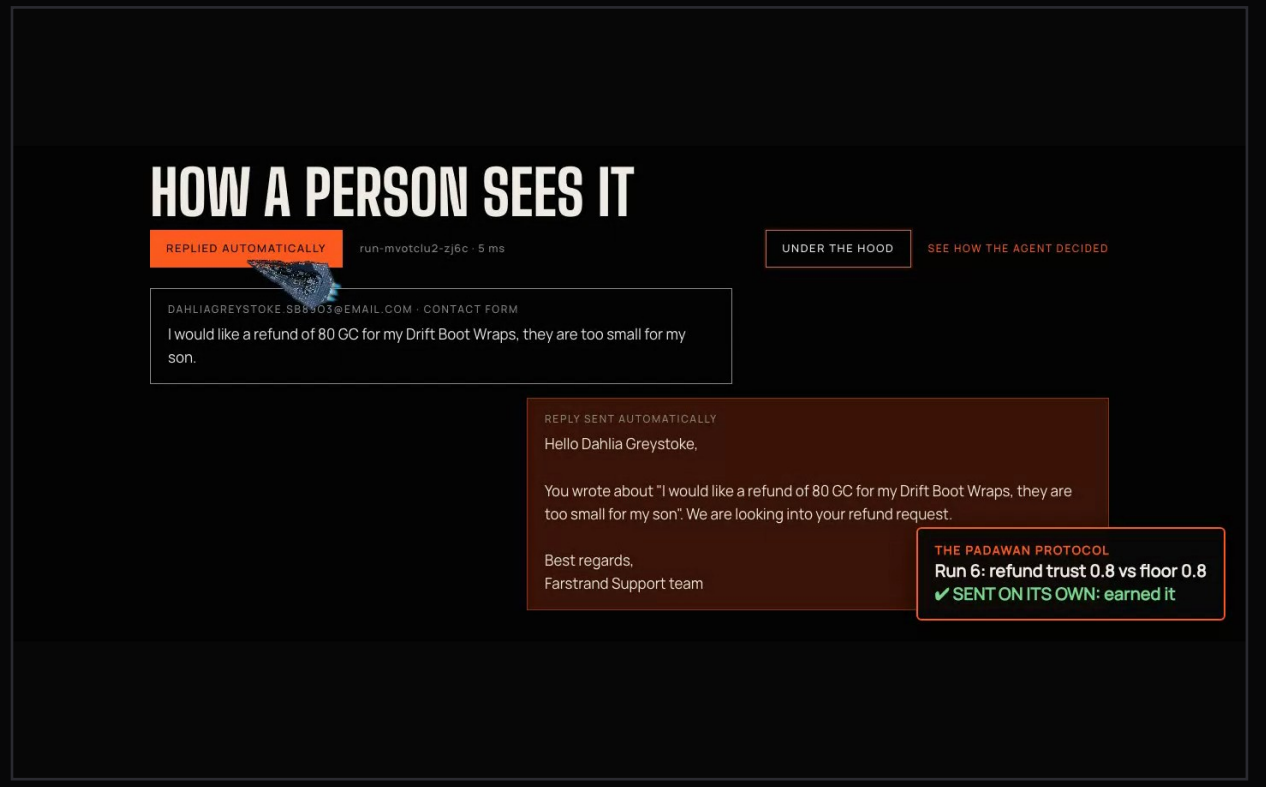

Recorded in sample mode (same maths as live, where Zetaris keeps the count)

1 HELD | 2 HELD | 3 HELD | 4 HELD | 5 HELD → Affirm → **Run 6** 0.80 FLOOR · SENDS ON ITS OWN → 1 Reject → 0.75 → **Run 7** HELD AGAIN

Runs 1–5 held · trust 0.55 → 0.75

A small stack with no runtime dependencies

MODELS	Nemotron 3.5 Lightning via Nightshift · DiffusionGemma via OpenJev
ENGINES	Meterless: Scout Intent, H-MEM, World Model
DATA	Zetaris: JEDI_PACK raw layer + Unified Semantic Layer
RUNTIME	Node 22, standard library only · Docker node:22-slim

874
TESTS

873
PASS

0
FAIL

1
SKIPPED

Built with Claude Code (through Glitch), orchestrated in Orca. SAMPLE_MODE=true runs the whole loop with no keys.

Next: make the simulated parts real

01 Teams escalation

The API moves the ticket, not only the page

02 Sign-off

Send the email and record the signature

03 Fallback proof

Record the local-lane fallback run

04 Impact

Measure time saved and wrong replies avoided

05 Real email and review intake

Today only the contact form is live; email and Trustpilot are simulated

06 Risk-tiered legal routing

Legal only for medium/high-risk threats, earned through the same trust ledger, a person always signs

TRY IT LIVE

jedi.blastworks.ai/shop/contact.html
then the staff desk at /shop/desk.html · live until 6 Nov